ml-journal

Persistent Session Audit Trail for Claude Code

Every decision, issue, and experiment — captured in a machine-queryable log that survives compaction and session boundaries.

A Claude Code plugin — install once, use in any git repo

The Problem with Ephemeral Sessions

Claude Code sessions are stateless. When the context compacts or a new session starts, the reasoning trail disappears.

- **Decisions are made but not recorded** — future sessions re-litigate the same tradeoffs
- **Issues discovered in one session go untracked** — they resurface as surprises later
- **Post-mortems happen from memory** — incomplete, biased toward recent events
- **Handoff between sessions requires manual summaries** — which are never written

> **Result:** Teams run the same failed experiment twice, lose the rationale behind architectural choices, and can't reconstruct what actually happened during an investigation.

What ml-journal Is

A Claude Code plugin that maintains a **structured, append-only JSONL audit trail** for every session — with ten skills for logging, querying, and synthesizing the record.

Layer	Component	Role
Judgment	Claude skills	Extract, classify, construct log arguments from conversation
Mechanical	Python scripts	Validate fields, serialize, append to journal
Storage	`.project-log/journal.jsonl`	Per-repo, append-only, machine-queryable

No background daemons. No external dependencies beyond `python3` and `git`.

> **Install:** `/plugin install claude-ml-journal@ml-debate-lab` then `/log-init` in any git repo.

Ten Skills, Two Jobs

Skills split cleanly into **write** (add to the journal) and **read** (query and synthesize).

Skill	Job	Direction
`/log-entry`	Log a typed entry from conversation context	Write
`/log-commit`	Git commit + journal log in one step	Write
`/checkpoint`	Save session state for handoff or post-compact recovery	Write
`/resume`	Load and display the most recent checkpoint	Read
`/log-status`	Quick overview — counts, unresolved issues, recent commits	Read
`/log-list`	List entries by type with optional time filter	Read
`/log-summarize`	Prose synthesis of all entries of a given type	Read
`/log-init`	One-time repo setup	Setup
`/research-note`	Session/day-scoped formatted note — shareable, PR-ready	Synthesis
`/research-report`	Synthesize `RESEARCH_REPORT.md` from full journal + git history	Synthesis

Invoking Skills — Two Ways

Skills can be invoked explicitly with a slash command, or in plain conversation.

Slash command:

```
You: /checkpoint
You: /log-summarize decisions
You: /research-note
```

Natural language — Claude recognises the intent:

```
You: save our progress before I compact
You: what decisions have we made so far?
You: write up today's work as a shareable note
```

Both routes invoke the same skill. Slash commands are convenient shortcuts — they're not required.

Entry Types

ml-journal enforces a fixed schema per entry type. Claude infers the type from conversation context — you don't have to label it.

Type	Required Fields	Confirm?
`issue`	description, severity	No
`resolution`	description	No
`decision`	description, rationale	Yes
`discovery`	description	Yes
`hypothesis`	description	Yes
`experiment`	description, verdict	Yes
`post_mortem`	description, what_failed, root_cause	Yes
`memo`	description	No
`checkpoint`	in_progress	Yes
`git`	commit_hash, message, branch	Yes (via `/log-commit`)

Medium-stakes types (decisions, experiments, post-mortems, discoveries, hypotheses) require confirmation before writing.

Logging an Entry

`/log-entry` reads the conversation, infers type and fields, and writes to the journal.

Example — logging a decision:

```
You: We decided to use cosine similarity over Euclidean distance because
     the embedding magnitudes vary by document length. Log this.

ml-journal: decision | "Use cosine similarity for document embeddings"
            Rationale: embedding magnitudes vary by document length,
                       making Euclidean distance unreliable.
            ► Log this entry? (y/n)
```

What gets written:

```
{"id": "a3f2...", "type": "decision", "timestamp": "2026-04-09T06:00:00Z",
 "description": "Use cosine similarity for document embeddings",
 "rationale": "embedding magnitudes vary by document length"}
```

Proactive Logging via CLAUDE.md

`/log-init` optionally injects logging rules into your repo's `CLAUDE.md`:

```
► Add proactive journal logging rules to your CLAUDE.md? (y/n)
```

If accepted, Claude auto-proposes entries at natural pauses — without being asked:

Pattern	Proposed type
User confirms a direction	`decision`
Unexpected finding changes approach	`discovery`
Bug identified and explained	`issue`
Fix confirmed working	`resolution`
Root cause understood	`lesson`
Experiment verdict clear	`experiment`

Rules: one proposal per event — not mid-investigation. Never re-proposes if declined. Chains `issue → resolution → lesson` at completion, not as three interruptions.

> Without CLAUDE.md rules, logging is always explicit. With them, Claude surfaces the right entry type at the right moment.

Committing and Logging in One Step

`/log-commit` stages, commits, and writes a `git` journal entry — all from a single command.

```
You: /log-commit
```

ml-journal:

1. Runs `git status` and `git diff --stat HEAD` — shows you what will be committed
2. Synthesizes a commit message from conversation context
3. Proposes a `git` journal entry with `diff_summary`
4. Asks for confirmation
5. Commits, captures the hash, logs to journal

> **Why one step?** Separating commit from log breaks the audit trail. The journal entry captures the *why* (diff_summary) that the commit message doesn't have room for — but only if they happen together.

Session Handoff: Checkpoint and Resume

The two highest-value skills for multi-session work.

`/checkpoint` — saves current session state before a `/compact` or end of session:

- What's in progress
- What's blocked and why
- What the next action is
- Key decisions made this session

`/resume` — loads the latest checkpoint at the start of a new session:

- Surfaces the in-progress state directly into context
- No manual summary required
- Works across compactions — the checkpoint survives even if the conversation doesn't

> **Pattern:** Run `/checkpoint` before every `/compact`. Run `/resume` at the start of every session. The journal becomes the session memory that Claude Code doesn't have natively.

Querying the Record

Three read skills for different levels of detail.

`/log-status` — quick dashboard:

```
Last checkpoint: 2026-04-09 06:00 | "Finishing Phase 3 scoring"
Entries this session: 4 decisions, 2 issues, 1 experiment
Unresolved issues: 2
Recent commits: fix: scorer bug in phase_score.py
```

`/log-list issue --since 7d` — enumerate recent entries of a type

`/log-summarize decision` — prose synthesis across all decisions:

> *"Three architectural decisions shaped the pipeline: cosine similarity was chosen over Euclidean distance due to magnitude variance, concurrent execution was adopted after sequential execution proved too slow, and the validation gate was tightened after two false-positive smoke tests."*

Research Note vs. Research Report

Two synthesis skills for different audiences and scopes.

`/research-note` — session or day scoped, shareable

- Draws from recent journal entries (or an existing `log-summarize` output)
- Produces `RESEARCH_NOTE_<date>.md` — 40–80 lines
- Sections: Summary, Key Decisions, Discoveries & Results, Issues, Current State, Next Steps
- Use after a work session, before a PR, or to share a daily update

`/research-report` — full project or phase scope

- Dispatches `report-drafter` subagent to isolate heavy context ingestion from the main session
- Draws from all journal entries, full `git log`, and supplementary markdown files
- Produces `RESEARCH_REPORT.md` — comprehensive retrospective
- Sections: Problem Statement, Timeline, What Was Tried/Failed/Worked, Key Decisions, Issues and Resolutions, Open Questions
- Use at end of phase or project, or to onboard a new collaborator

Chain Workflows — Bug Fix & Investigation

Proactive logging rules chain entry types together. Claude proposes each step at the right moment.

Bug fix — three entries, one event chain:

bug fix chain

Investigation — hypothesis → experiment → outcome:

investigation chain

Chain Workflows — Session & Synthesis

Session boundary — state survives compaction:

session chain

Synthesis ladder — from raw entries to shareable artifacts:

synthesis chain

Optional Hooks

Two hook scripts enable automatic checkpoint/resume without manual invocation.

Hook	Script	Behavior
`PreCompact`	`journal-precompact.sh`	Auto-writes a checkpoint before `/compact`
`SessionStart`	`journal-session-start.sh`	Injects latest checkpoint as session context

To enable, add to `.claude/settings.local.json` (per-machine, gitignored):

```
{
  "hooks": {
    "PreCompact": [{ "type": "command",
      "command": "bash plugins/ml-journal/journal-precompact.sh" }],
    "SessionStart": [{ "type": "command",
      "command": "bash plugins/ml-journal/journal-session-start.sh" }]
  }
}
```

Hooks are not registered by default — install only if your environment permits them.

When to Use ml-journal

Use ml-journal when:

- You run multi-session investigations and need context to survive compaction
- You make architectural decisions that future sessions will need to understand
- You want a reconstructable audit trail for post-mortems or research narratives
- You're using ml-lab and want decisions and experiment outcomes logged alongside git history

You get the most value from:

- `/log-commit` as a commit habit — git + journal in every commit
- `/checkpoint` before every `/compact`
- `/resume` at the start of every session
- `/log-summarize` before a post-mortem to reconstruct what happened

> **Works well with ml-lab:** ml-lab generates decisions, discoveries, and experiment outcomes — ml-journal captures them. Run both plugins together for a complete investigation record.